

A Virtual Middleboxes Placement Algorithm in Multi-tenant Datacenter Networks

Jian Wang^a, Xuewei Zhang^a, Xiaoliang Wang^{a,*}, Sanglu Lu^a

^a*National Key Laboratory for Novel Software Technology, Nanjing University, Nanjing, P.R. China*

Abstract

Hardware middleboxes are widely used in current cloud datacenter to provide network functions such as firewalls, intrusion detection system, load balancers, etc. Unfortunately, they are expensive and unable to offer customized functions for individual tenant. To overcome this issue, there is an increasing interest in deploying software middleboxes to enable flexible security, network access functionality. This paper addresses the software middleboxes placement problem with minimum bandwidth guarantee. We first specify the model of tenants' requirement that specifies the need for virtual machines of application and middleboxes, as well as communication traffic. A virtual middlebox placement algorithm called EMB is then proposed to offer predictable network performance for each accepted tenant, and minimize datacenter bandwidth utilization. Extensive simulation results based on current large-scale datacenter networks verify that EMB is effective and provides network performance guarantee for tenants.

Keywords: **network virtualization**, middlebox, VM placement, bandwidth guarantee, datacenter

1. Introduction

Large public clouds, such as Amazon EC2 and Windows Azure, provide infrastructure as a service (IaaS) for enterprises and individuals. Besides the

*Corresponding author

Email address: waxili@nju.edu.cn (Xiaoliang Wang)

ability to provide computing resources, public clouds are expected to support
 5 network functionality, such as load balancers, firewall, intrusion detection system, etc. Currently, network functions are widely deployed by the expensive hardware middleboxes (HW MBs). As illustrated in Fig. 1, hardware middleboxes are placed at the point close to the data center switches. Since hardware middleboxes are shared by all tenants, it is hard to manage these expensive
 10 devices to provide customized functions for each tenant. Also, heavy congestion may happen at links around those middleboxes. To address these issues, cloud providers and researchers introduce software middleboxes (or, interchangeably, virtual middleboxes, middlebox VMs) in the cloud [1, 2].

Providing virtual middleboxes with guaranteed performance is challenging
 15 for cloud service providers due to several reasons. First, we need to model various requirements of tenants, which consist of application VMs, middlebox VMs as well as the connecting network. Second, the performance of software MBs is not easy to predict in real environment due to many factors like the number of rules, algorithms and implementations. The experience based selection of
 20 MBs may under- or over-estimate resource consumption. Last but not least, it requires an algorithm to efficiently place virtual middleboxes in cloud service infrastructure.

Inspired by recent works on application-driven bandwidth guarantee [3, 4], we consider the middlebox placement strategy to overcome these challenges.
 25 Tenant's requirements is represented by a weighted directed acyclic graph (DAG). The nodes consist of diverse virtual middleboxes and a set of homogeneous application VMs, both of which are dedicated for each tenant to process arrival traffic. The edge of DAG indicates the communication request of middlebox-to-middlebox, middlebox-to-application traffic and intra-applications. The weight
 30 on the edge is the minimum communication requirement between application VMs [2, 5]. The inter-tenant communication [3] is also structured by a directed network. The arrival traffic from other tenants will be forwarded to the MBs of the target tenant before arriving at its application VMs.

Based on the model of tenant request, we design an effective algorithm that




Figure 1: Datacenter network and hardware middlebox deployment

carefully places the required MBs and application in form of VMs in datacenter to maximize the utilization of resources (e.g., computation and network). To address the bandwidth constraint in an oversubscribed datacenter networks, we adopt the depth-first search (DFS) allocation algorithm in the tree-like datacenter topology to minimize the network bandwidth occupancy for each accepted tenant. Moreover, to make good use of both physical machine resource and network bandwidth resource, we consider dividing the middleboxes and the application VMs proportionally and putting them in the same server or rack so as to occupy as many virtual machine slots as possible to eliminate unnecessary communication cross racks due to the unbalanced assignment of traffic processing capacity. Further, we try to colocate more application VMs to reduce the intra-tenant communication cost through carefully exchanging the positions of middleboxes and application VMs. If free slots or bandwidth is not enough, the middlebox chain and application VMs have to be placed separately but in closed sub-tree to reduce the bandwidth consumption caused by communication intra- and inter-tenants.

To the best of our knowledge, we are the first to model the virtual middlebox and application network placement problem in multi-tenant datacenter networks with minimal performance guarantees. Extensive simulations on large-scale datacenter topology have shown the effectiveness of our VMs placement algorithm. The accept ratio approaches to 90% with regard to a variety of tenants' requests in a large-scale datacenter topology. Compared with the existing

placement algorithm, our algorithm can also greatly reduce the bandwidth utilization of high level links while making the best of low-level links' bandwidth resource. Compared with the hardware middlebox placements in current data-
60 centers, our VMs placement shows significant saving on network bandwidth but provides the performance guarantee. Generally, the network function virtualization enables the cloud services provider to offer customized services with lower cost and our study will be helpful to build more flexible software-based cloud services to meet diverse requirements of the tenants in the datacenter network.

65 The rest of this paper is organized as follows. Sec. 2 reviews the background and related works on tenant's performance guarantee in the multi-tenant datacenter. Sec. 3 describes our abstraction of middlebox and the communication model. The VMs placement algorithm is explained in Sec. 4, and the simulation based performance evaluation is demonstrated in Sec. 5. Sec. 6 concludes this
70 paper.

2. Background and Related Works

Multi-tenant datacenter Large-scale datacenter is the core infrastructure for cloud based services. Datacenter network is organized in a multi-level tree-like topology. See Fig. 1, each rack typically contains 40+ servers, each of
75 which is configured a 10Gbps network interface card (NIC) connecting to the top-of-the-rack (ToR) switch. The ToR switch will balance the traffic to multiple aggregation switches. All aggregation switches are further connected to the core switches. Note that the communication bandwidth between adjacent layers is oversubscribed by a ratio from 2:1 to 10:1, which may vary slightly with different
80 configurations [3, 6].

In multi-tenant data center environments, a tenant can rent multiple virtual machines (VMs) on physical servers. The VMs are configured with different amounts of CPUs, memories and storage resources. From the network perspective, we usually model the computation and storage resource to be the number
85 of slots requested in physical machine. All VMs assigned to a tenant are inter-

connected to form a virtual network (VN) such that VMs can communicate with each other to exchange data. The traffic inside a multi-tenant datacenter carries external traffic to and from Internet, intra-tenant traffic between a tenant’s application VMs and inter-tenant traffic. Specifically, the inter-tenant traffic
90 consists of flows between tenants and flow between tenant and providers [3], which can also be considered as a special tenant. Inter-tenant communication increases with the increase of services offered by cloud providers. Based on recent measurement in public cloud provider, 65% of traffic happens between VMs belonging to one tenant VN and 35% is among different VNs or tenant-provider
95 communication [3, 7].

Middleboxes in multi-tenant datacenter Network functions are also necessary and provided in the form of hardware middleboxes in current datacenters. However, hardware middleboxes are known to have several drawbacks. For example, working as a shared resource, the hardware middleboxes can only
100 provide common services used by all tenants instead of providing customized tenant requests. Once the middlebox failed, all the tenants will be affected. Besides, the fixed locations of hardware middleboxes result in aggregated network bandwidth consumption and cause network congestion [8, 9].

An alternative solution to hardware middleboxes is network function virtualization (NFV) [10]. Deploying software middleboxes will greatly improve the
105 flexibility, manageability and cost-efficiency. Software middlebox is not necessary to be shared and able to provide customized services with performance guarantee for each dedicated tenant. It has been shown that the software middleboxes are able to offer equivalent functionalities as the corresponding hardware
110 implementation [11, 12, 1, 8]. The advance of software defined networking (SDN) further facilitates the deployment and management of software middlebox running in datacenters. Although there remains unsolved challenges before it can be integrated into the cloud ecosystem, the perspective of deploying flexible software middlebox for isolated tenants has motivated us to extend the previous
115 works to realize predictable middlebox virtual networks through effective virtual middleboxes placement.

Placement of services with performance guarantee Since the underlying network is shared by tenants, the application can be affected by variable and unpredictable network performance. A series of works have been proposed to enable traffic isolate and provide bandwidth guarantee, including Oktopus [5], FairCloud [13], etc. The above works focus only on intra-tenant communication. The inter-tenant traffic isolation is addressed by Hadrian [3]. Through hierarchical hose model, the communication dependencies can be explicitly declared. The cloud provider then allocates resources according to the specification of tenants.

Most of previous works only considered the bandwidth guarantee for applications. As far as we know, little work addresses the bandwidth guarantee for virtual MBs in multi-tenant datacenters [14]. The NFV cloud computing platforms, like NFV OPEN LAB [15] and CloudNFV [16], provide NFV management and orchestration as applications. Theoretically, to obtain effect use of resources, various NFV placement objectives have been studied, for example, remaining data rate, latency, number of used network node as well as distance and setup cost, see [17, 18] and the references therein. However, they could not provide bandwidth guarantee for tenants. Stratos [2] considered the middlebox placement in multi-tenant datacenters through a systematic way by avoiding network congestion. In contrast, we first study the bandwidth consumption problem of virtual middleboxes and then propose effective placement algorithm to make good use of network resources.

3. Abstraction of Virtual Middlebox Network

In a multi-tenant datacenter, by leveraging the network function virtualization (NFV), cloud providers provide middlebox as a fundamental network service to tenants, like both computing and storage resources. A tenant could specify their middlebox and application (i.e., computation/storage) requirements in forms of VMs and the expected minimum bandwidth requirement for intra- and inter-tenant communication. Cloud providers will then assign VMs for middle-




Figure 2: Intra-tenant communication and inter-tenant communication

boxes (in short, MB VMs) and applications (in short, APP VMs) according to the tenant's request. The cloud providers also enforce the middlebox policies, i.e., the sequence of middleboxes that a flow must traverse. In this section, we explain the model for APP VMs and MB VMs. We then introduce the
150 abstraction of virtual networks associated to one tenant.

3.1. Model of APP VMs

To provide predictable network performance for the communication of N APP VMs in one tenant, users typically specify the minimum bisection bandwidth B^{in} for individual APP VM. Analogous to previous proposals [3, 13], the
155 requirement of the minimum bandwidth guarantee for each APP VM belonging to the same tenant can be abstracted by the *hose model*. In the hose model (see the thick black arrows in Fig. 2, tenant P), the application VMs are connected to a logical switch by a link of capability B^{in} .

Similarly, for the cross tenants communication, a tenant can specify the
160 bandwidth request for a single APP VM. The minimum bandwidth guarantee for inter-tenant traffic of each APP VM is defined by B^{ex} . We extend the VN model of Hadrian [3] to support multiple external links. If the current tenant communicates to multiple tenants. The minimum bandwidth for each communication pair is guaranteed to be B^{ex} . The communication pattern between
165 tenants can be expressed by a directed graph. We use the definition of *communication dependency* in [3] to explicitly specify the peer-to-peer traffic pattern. A tenant with a dependency of $\{*\}$ means that he could allow any other tenants' communication request. We call this kind of tenant open tenant while the ten-

ant who has communication requests to open tenants are called client tenant.

170 When the dependency is set as $P:\{U, V\}$, it means that the communication from tenant U and tenant V to tenant P are allowed. An example is illustrated in Fig. 2, the thin blue arrows show the communication between tenant U and tenant P, as well as tenant V and tenant P.

3.2. Model of MB VMs

175 Middleboxes perform different network functions to meet traffic demand of applications. A tenant may vary network functions simultaneously to process packets of different traffic class, as shown in Fig. 3. Each request consists of a sequence of various MBs, such as firewall (FW), intrusion detection system (IDS), redundancy elimination(RE), load balancer (LB) and NAT Proxy, etc., 180 which provide functionality of access control and traffic engineering.

Generally, the communication between MBs can be abstracted by a directed-acyclic graph (DAG), where the vertex represents the type of middleboxes and the arc indicates the traffic flow between them. The external traffic should pass through the specified middleboxes one by one according to the DAG. The sequence can not be violated. Otherwise, the expected functionality will be ruined. 185 In addition, we assume that the traffic still keeps the same when multiple MBs are introduced. This is reasonable because the traffic is the minimum bandwidth guaranteed for both intra- and inter-tenant communication. The normal traffic could pass all the rule checking and will not be blocked by middleboxes like FW or IDS. For those middleboxes, e.g., RE or LB, which change the traffic 190 load, they are usually placed at the end of the DAG. The variation of traffic is specified by the proposed parameters of B^{in} and B^{ex} in the aforementioned APP VMs model.

Since a tenant may need to apply different middlebox policies to different traffic class, the tenant request can be classified into multiple MBs chains associated with different applications [2, 17, 14]. See Fig. 3 as an example. The Internet traffic will pass Firewall, Intrusion Detection System and then NAT proxy, while the traffic from other tenant may only pass Firewall and Load Balancer.




Figure 3: Abstraction of MBs: Two MBs chains

Therefore, a middlebox chain model can be defined by a tuple $\langle T, M \rangle$, where
 200 the type of MBs is denoted by a vector T , and the number of required MB VMs
 is denoted by a vector M . The direction of traffic flow is implicitly expressed by
 the sequence of items in the vector T . For instance, the first MB chain in Fig.
 3 can be expressed as $\langle T, M \rangle$, where $T = \{\text{FW}, \text{IDS}, \text{Proxy}\}$, $M = \{2, 4, 2\}$.
 It means that the tenant request for 2, 4 and 2 VMs of FW, IDS, and Proxy
 205 respectively. The traffic will arrive at FW first, and then go to IDS and pass
 Proxy before reaching the application VMs.

3.3. Virtual Network Model for MB VMs and APP VMs

We combine the MB model and APP model together to form the virtual
 network (VN) model to represent the tenant's request. The interaction between
 210 VMs of the same tenant is expressed by the hose model. The middlebox chain
 should be attached to the ingress path of the external link of the VN. Such that,
 all external traffic coming into APP VMs must traverse through these specified
 middlebox chains which provide access control and other network functions.
 CloudMirror [4] provided a new network abstraction that allows the applications
 215 to specify their directional and specific communication pattern. We adopt a
 cascade of *Tenant Application Graph (TAG)* proposed in [4] to specify the VN
 model.

Fig. 4 (a) shows an example, where a sequence of MBs is attached at the
 external link of APP VMs of tenants. The arrow shows the traffic flow. Ar-
 220 rival traffic of tenant A will pass through all these middleboxes before being
 distributed to the application services. Similarly, the departure traffic of tenant
 A is forwarded to its destination (tenant P) which needs to pass through the




Figure 4: Virtual Network Model. (a) the communication between two tenants; (b) cascade of TAG

middleboxes of tenant P first ¹. Fig. 4 (b) uses the cascade of TAG model to abstract the communication dependency of VNs. The external traffic passing through each type of middleboxes in sequence are expressed by the dedicated directional arrows. The internal traffic of APP VMs is expressed by bi-directional arrows, each of which has a capacity of B_A^{in} (tenant A) or B_P^{in} (tenant P). The external traffic of each VM of tenant A (resp. tenant P) is bounded by B_A^{ex} (resp. B_P^{ex}).

In general, two kinds of traffic may go through the MBs before come to the Application VMs, 1) the Internet traffic; 2) the traffic from other tenant. When declaring a connectivity relationship, this external link can only communicate with those specified VNs. The unauthorized traffic will be blocked by MBs. Based on the models of APP VMs and MB VMs, a tenant can specify the connectivity of an external link which receives (or sends) traffic from any VN. In summary, a tenant's request can be specified by multiple virtual networks (VNs) and each VN can be specified by a tuple $\langle N, B^{in}, B^{ex}, dependencies, T, M \rangle$.

Remark It is a remarkable fact that the behavior of current software middleboxes is actually not easy to predict in the real system. The performance depends on multiple factors, including workload like traffic demands, type and number of rules, and middlebox design, e.g., algorithm design, systematical im-

¹Some corporations may have the requirement to detect the departure traffic through additional middleboxes. We leave this for future work.

plementation, etc. Therefore, the scalability of MBs is important and need to be considered by the cloud service provider. Once the virtual network are placed in datacenter, it is hard to find a proper resource to migrate MBs and re-allocate live flows across MBs with performance guarantee [1]. We discuss the scalability issue in Appendix.

4. Virtual Middlbox Network(VMN) Deployment

Based on the tenant’s request, the cloud provider allocates enough resources for MBs and APPs. The VMs request will be satisfied by providing enough free slots on physical machines (PMs). The network function demand consists not only the bandwidth request but also the capability to process a specific class of traffic in order. We present a VM placement algorithm that bridges the gap between the high-level virtual middlbebox network model to the low-level physical network. The placement strategy is realized by a logically centralized controller, which manages datacenter operations through a global view of the network topology. SDN based traffic engineering manipulates the underlying network routers to explicitly direct the class of traffic through each middlebox sequentially. If there are not sufficient resources, the request will be rejected.

The virtual machine placement problems have attracted a lot of attentions recently, which are often formalized as a multi-dimensional packing problem with various constraints. This problem is NP-hard [19, 20]. Therefore, we search for effective heuristic solution to maximize the datacenter resources utilization. Before introducing the VMs placement algorithm, we first explain some practical constraints in our design. Next, we show the bandwidth requirement for the tenant request. Finally, we present our placement strategy by combining two simple placement algorithms.

4.1. Bandwidth Constraint within a Physical Machine

Existing operating systems typically copy the data from user space to kernel space and then move the data to the buffer of NIC. The data copies cause a long




Figure 5: Bandwidth constraint of NIC in PM

270 delay during packet transmission. Recent works allows applications to access
 the NIC card directly without involving kernel processing which direct DMA
 to and from a VM, like Single Root I/O Virtualization (SR-IOV) [21]. Intel
 SR-IOV allows multiple VMs to access the NIC through L2 address. Most data
 center networks have enable the SR-IOV in their system to accelerate the data
 275 processing². When using SR-IOV, packets are switched on a per-port basis in
 the NIC between VMs on a shared port, as that shown in Fig. 5. The hypervisor
 configures the NIC switch on the SR-IOV network adaptor to provide network
 connectivity to the VMs. However, the communication between VMs is bounded
 by the NIC switch capacity. NetVM [22] tried to overcome this limitation by
 280 providing memory sharing without passing through the NIC. However, due to
 some problems like the potential security thread in multi-tenant datacenter, this
 new feature is not applied in current commodity environment yet. In general,
 the bandwidth constraint inside PM can be expressed that the traffic between
 VMs of the same PM is bounded by a maximum bandwidth, see the model at
 285 right side of Fig. 5. We need to take this constraint into consideration during
 VMs placement algorithm design.

²Virtual Ethernet Port Aggregator (VEPA) forwards packets out of the host and come
 back via an external switch to be transmitted to a VM. Since VEPA will cost more network
 resource, we do not consider this situation in this paper.




Figure 6: Bandwidth requirement

4.2. Bandwidth Required by VMN

Giving a virtual network request, we need to provide enough middleboxes to process the traffic from the external Internet or other tenants. We assume each instance of MB_k has a maximal traffic processing capacity c_k , So the minimum required number of MB_k is:

$$M_k = \left\lceil \frac{N * B^{ex}}{c_k} \right\rceil. \quad (1)$$

Then we can obtain the total number of VMs required for the middleboxes and applications.


We again assume that the total traffic is shared equally among each instance of MB_k . And the practical required bandwidth b_k for each instance can be calculated by :

$$b_k = \frac{N * B^{ex}}{M_k}. \quad (2)$$

290 Based on above calculations, we analyze the bandwidth consumption to achieve a better understanding of the MBs placement problem. Assume that we have found two candidate subtrees that have enough slots to hold the required VMs as shown in Fig. 6. The total number of instances required for MB_k is indicated by M_k . And m_k is the number of instances placed for MB_k



(a) TRAIN: dividing virtual middleboxes network horizontally



(b) TRUCK: dividing virtual middleboxes network vertically.

Figure 7: Two partition strategies to distribute the VMs among subtrees

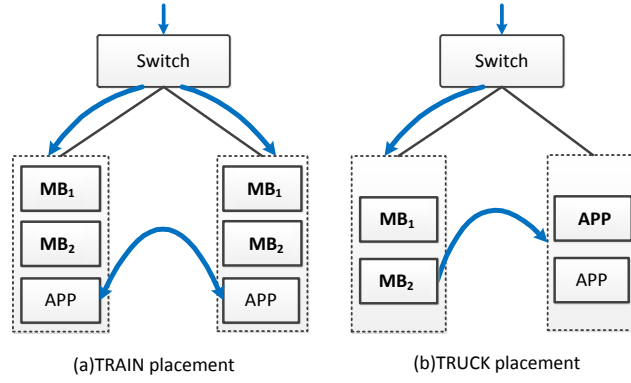


Figure 8: Illustration of placement algorithms

295 in the left subtree. Correspondingly, the number of instances for MB_k in the right subtree is $M_k - m_k$. The n_{app} and $N - n_{app}$ are the number of APP VMs in the two subtrees respectively. Suppose the arrival traffic is $N * B^{ex}$ and the middlebox instances prefer to send the traffic to its downstream middlebox instances inside the same subtree.

We then analyze the bandwidth cost of this situation. The traffic flows across and inside the subtrees are illustrated in thick and fine arrows respectively. Specifically, we find that the bandwidth consumption across the subtrees consists of four parts indicated by B_1 , B_2 , B_3 and B_4 . The B_1 is the external traffic (represented by the black arrows) from the switch to MB_1 which is equal to $N * B^{ex}$. The traffic that cannot be fully processed by MB_2 will be partly sent to the instances of MB_2 in the opposite subtree as shown in the blue arrow labeled by B_2 . Similarly, the traffic from MB_2 in the right to APP will be partly sent to the APP VMs in the left subtree (red arrow labeled by B_3) if the capacity of APP VMs in the right cannot serve the whole traffic from MB_2 in the same side. Besides, the bandwidth requirement (dotted arrow labeled by B_4) for intra-tenant communication is $2\min(n_{app}, N - n_{app}) * B^{in}$. For simplicity, we only consider the downlink bandwidth consumption since the uplink bandwidth has the similar characteristics. Based on above analysis, we can easily get the total downlink bandwidth requirement across the subtrees:

$$\begin{aligned}
bw &= B_1 + B_2 + B_3 + B_4 \\
&= N * B^{ex} + (m_1 * b_1 - m_2 * b_2) \\
&\quad + ((M_2 - m_2) * b_2 - (N - n_{app}) * B^{ex}) \\
&\quad + 2\min(n_{app}, N - n_{app}) * B^{in}.
\end{aligned} \tag{3}$$

For the general situation, by adding all the inter-middlebox traffic going

across the subtrees, we can get the following new equation:

$$\begin{aligned}
bw &= B_{switch,mb_1} + B_{mb_k,mb_{k+1}} + B_{mb,app} + B_{app} \\
&= N * B^{ex} + \sum_{k=1}^{n-1} |m_k * b_k - m_{k+1} * b_{k+1}| \\
&\quad + |m_n * b_n - n_{app} * B^{ex}| \\
&\quad + 2min(n_{app}, N - n_{app}) * B^{in}.
\end{aligned} \tag{4}$$

300 B_{switch,mb_1} is the incoming traffic for MB_1 from the switch which is a constant. $B_{mb_k,mb_{k+1}}$ and $B_{mb,app}$ are the incoming traffic for MB_{k+1} and APP VMs across the subtrees from MB_k respectively. B_{app} represents the intra-tenant communication bandwidth requirement.

4.3. VM Placement Strategies

305 Based on above bandwidth calculation, we can find two simple strategies to distribute the VMs among the subtrees. One approach is to always place all kinds of MB VMs in the same subtree proportionally with corresponding APP VMs. Each partition contains sufficient MBs to serve the APPs which can works independently. We call such approach 'TRAIN', since it works like a train
310 consisting of multiple cars of MBs. As shown in Fig. 7(a), it consists of three groups horizontally, each of which can process all arrival traffic to the APP VMs in the same group. Another approach is to make the same kind of MBs as a group and place them together in the subtree. We call this approach 'TRUCK' since it works like a truck to take cargo of the same kind. As shown in Fig.
315 7(b), MB_1 , MB_2 and APP are divided first and put to subtrees separately.

Fig. 8 shows the placement examples for above two simple strategies. The bandwidth consumption across the subtrees is represented by the blue arrows. Obviously, the placement strategy TRAIN can minimize the bandwidth requirement from MB_k to MB_{k+1} and APP, but it may have a large bandwidth demand for intra-tenant communication. In contrast, TRUCK can save intra-tenant bandwidth greatly but having a big waste of bandwidth from MB_k to MB_{k+1} or APP.




Figure 9: Illustration of the bandwidth variation after exchange

Considering the advantages of the two strategies, we seek to obtain a more powerful placement algorithm to achieve bandwidth saving by combining them. More specifically, to reduce the traffic $B_{mb_k, mb_{k+1}}$ and $B_{mb, app}$ crossing the subtrees caused by the unbalanced capacity between the upstream and downstream MBs or APPs, we can place the MBs and APPs proportionally in the two subtrees according to their traffic processing capacity like TRAIN. In addition, as n_{app} increases from zero to N , the intra-tenant bandwidth B_{app} increases from zero to $(N/2) * B^{in}$ when $n_{app} = N/2$ and then decrease to zero. So through collocating more APP VMs in one subtree like TRUCK, we can greatly save the intra-tenant bandwidth B_{app} .

Thus, we get the following two placement principles: 1) placing the MBs and APPs proportionally according to their capacity initially and then 2) collocating as many APPs as possible through exchanging the positions of MBs and APPs in the two subtrees. While the exchanging may result in a change of bandwidth consumption due to the break of balance of MBs and APPs capacity. So we need to carefully select the MBs to exchange.

We next analyze the change of the bandwidth before and after the exchange of the MB and APP. Consider a simple example as shown in Fig. 9. The tenant request has been placed in two subtrees proportionally according to the

placement principle 1 showed in Fig. 9 (a). So the bandwidth requirements $B_{mb_k,mb_{k+1}}$ and $B_{mb,app}$ are almost zero since the traffic from MB_k can be fully processed by MB_{k+1} s or APPs in the same subtree. Note that the number of
 345 APPs in the left subtree is larger than the number in the right side. We then try to exchange the MB_2 in the left side and APP in the right side to reduce the traffic across the subtrees between APPs. The result of exchange is shown in Fig. 9 (b). The dashed arrows show the change of traffic after exchange. The label associated with each arrow is the value of traffic change. The long-dashed
 350 arrows show the increase of the bandwidth in the two links. While the short-dashed arrow represents the decrease of bandwidth across the subtrees. After the exchange, the MB_2 in the left side cannot totally handle the traffic from MB_1 in the same side. This leads to an extra flow from MB_1 in the left to MB_2 in the right side expressed in the blue long-dashed arrow. Similarly, the APPs
 355 in the left side get additional capacity to process the extra flow from MB_2 in the right side due to the exchange showed by the green long-dashed arrow. For the intra-tenant communication, we get a decrease of bandwidth represented by the short-dashed arrow through increasing the number of colocated APPs in one side.

Thus, if $2B^{in} > 2b_2 + B^{ex}$, it is a better choice to exchange, i.e., the total bandwidth consumption in the two links will be reduced if the decrease is larger than the increase of the bandwidth after the exchange. And, we can easily get the following exchange conditions for any MB_k using the same analyzing method as shown above:

$$\begin{aligned} 2B^{in} &> b_k + B^{ex}, k = 1; \\ 2B^{in} &> 2b_k + B^{ex}, k > 1. \end{aligned} \tag{5}$$

However, the total decrease of the bandwidth cannot ensure that both of the links have a reduction of bandwidth. For a better understanding, we assume $b_2 = 200$, $B^{ex} = 600$, and $B^{in} = 502$ (in Mbps). Explicitly, MB_2 satisfies the exchange condition mentioned above. Consider the simple example shown in Fig. 9. The right link get a bandwidth reduction computed as $B^{in} - b_2 =$

502 – 200 = 302 Mbps after exchange. While the left link have a bandwidth increase of $(B^{ex} + b_2) - B^{in} = 600 + 200 - 502 = 298$ Mbps. So the total bandwidth has been reduced. But the exchange may lead to an unbalanced utilization of the two links. And the tenant request will be rejected if the remain bandwidth is smaller than 298 Mbps in the left link before the exchange. To avoid this situation, we reconsider the condition of exchange to ensure that both links have bandwidth savings. As shown in Fig. 9, the two links both get an extra traffic flow which causes a bandwidth increase from the opposite side, respectively. In addition, there is a bandwidth reduction of B^{in} for intra-tenant communication in both of the links. Thus, the condition to have bandwidth savings for both links is given as

$$B^{in} - (B^{ex} + b_2) > 0 \quad \text{and} \quad B^{in} - b_2 > 0, \quad (6)$$

i.e., if $B^{in} > (B^{ex} + b_2)$, the bandwidth will be saved for both of the links. Similarly, we can easily obtain the following improved conditions to get a more proper exchanging for any middlebox MB_k .

$$\begin{aligned} B^{in} > b_k \quad \text{and} \quad B^{in} > B^{ex}, k = 1; \\ B^{in} > b_k + B^{ex}, k > 1. \end{aligned} \quad (7)$$

360 Thus, using the improved exchangeable conditions to select the proper type of middleboxes for exchanging³, we cannot achieve a total bandwidth saving but ensuring a bandwidth reduction for each of the links when placing the tenant request.

4.4. VM Placement Algorithm

365 Based on above placement principles and exchangeable conditions, we design an effective heuristic VM placement algorithm that achieves bandwidth savings while providing guarantees. The basic idea behind the algorithm is to place MBs

³Equ. 7 is a sufficient but not necessary condition for bandwidth saving through exchanging, and we can also have bandwidth saving even the VMs are placed randomly initially by exchanging using Equ. 7.

and APPs proportionally initially to minimize the bandwidth consumption due to the unbalanced capacity assignment between middleboxes and application
 370 VMs; and then try to colocate more APPs to reduce the bandwidth cost of intra-tenant communication through exchanging MBs and APPs.

Algorithm 1 depicts the process to handle the tenant request. The algorithm starts with *PlaceTenant* (line 1), which takes a tenant request t as an input. To save the precious core level bandwidth, we traverse the datacenter topology in a
 375 bottom-up manner to search for a valid lowest level subtree that has enough VM slots for t (line 2). For the candidate subtree, we place t in it with *AllocateVM* (line 4) and then check whether the bandwidth requirement can be satisfied (line 5). If the trial fails due to the lack of bandwidth, *DeAllocate* (line 7) is invoked to release the resources reserved for the tenant request. And the
 380 algorithm continues to find another subtree in the same level or upper level if there is no subtree in the same level (line 8).

AllocateVM (line 10) is a recursive function, which takes a tenant request t , its APP and MB VMs number and a subtree st as inputs. The function first checks if the st is the lowest level subtree or physical server: if so, it allocates
 385 corresponding number of slots for t according to the parameters of N_{app} and N_{mb} in the server and then returns. If st is a higher level switch, we then invoke *AllocateVM* recursively for each subtree v of st which has available VM slots. For the subtree v that cannot provide enough slots for the unallocated MB and APP VMs, we first attempt to place appropriate number of MBs and APPs
 390 according to the placement principle¹ to fill the subtree v and then colocate as many APPs as possible by exchanging MBs and APPs using the exchangeable conditions of Equ. 7. After that, we invoke the function *AllocateVM* (line 24) for the adjusted MB and APPs. While for the subtree of st that has sufficient slot resource, *AllocateVm* is invoked directly (line 26).


395 In Algorithm 2, *ExchangeMbApp* (line 1) is invoked to exchange the positions of MBs and APPs. The function first identifies the types of middleboxes that can achieve bandwidth saving through exchanging their and APPs positions by using the exchangeable conditions of Eq. 7 (line 2). If the exchangeable

Algorithm 1 VM Placement Algorithm


```
1: function PLACETENANT(Tenant t)
2:   st = FindlowestSubtree(t, null)
3:   while st do
4:     AllocateVM(t, t.mb, t.app, st)
5:     if ReserveBw(t) then
6:       return true
7:     DeAllocate(t)
8:     st = FindLowestSubtree(t, st)
9:   return false
10: function ALLOCATEVM(t,  $N_{mb}$ ,  $N_{app}$ , st)
11:   /* $N_{mb}$  stores the number of instances for each type of middleboxes, while
      $N_{app}$  is the number of application VMs.*/
12:   if level(st)==0 then /*st is a server*/
13:     PlaceMB(t,  $N_{mb}$ , st)
14:     PlaceAPP(t,  $N_{app}$ , st)
15:   else
16:     for each subtree v of st do
17:       if AvailableVM(v)>0 then
18:         if AvailableVM(v)<TotalVM( $N_{mb}, N_{app}$ ) then
19:           place  $N_{v\_mb}$ MBs and  $N_{v\_app}$ APPs
20:           proportionally under v.
21:            $N_{mb} = N_{mb} - N_{v\_mb}$ 
22:            $N_{app} = N_{app} - N_{v\_app}$ 
23:           ExchangeMbApp( $N_{mb}, N_{app}, N_{v\_mb}, N_{v\_app}$ )
24:           AllocateVM(t,  $N_{v\_mb}$ ,  $N_{v\_app}$ , v)
25:         else
26:           AllocateVM(t,  $N_{mb}$ ,  $N_{app}$ , v)
```

Algorithm 2 Exchange Algorithm

```
1: function EXCHANGEMBAPP( $N_{mb}, N_{app}, N_{v\_mb}, N_{v\_app}$ )
2:    $mb = \text{GetExchangeableMB}(t)$ 
3:   if  $mb \neq \text{null}$  then
4:      $num$  is assigned as the number of types
5:     of the exchangeable MBs.
6:      $k = 0$ 
7:     if  $N_{app} > N_{v\_app}$  then
8:       while  $k < num$  and  $N_{v\_app} \neq 0$  do
9:         if  $N_{mb}[mb[k]] == 0$  then
10:            $k++$ 
11:           continue
12:            $N_{mb}[mb[k]] - -$ 
13:            $N_{app}++$ 
14:            $N_{v\_mb}[mb[k]]++$ 
15:            $N_{v\_app} - -$ 
16:       else
17:         while  $k < num$  and  $N_{app} \neq 0$  do
18:           if  $N_{v\_mb}[mb[k]] == 0$  then
19:              $k++$ 
20:             continue
21:              $N_{v\_mb}[mb[k]] - -$ 
22:              $N_{v\_app}++$ 
23:              $N_{mb}[mb[k]]++$ 
24:              $N_{app} - -$ 
```




(a) Accept rate under different B^{in}




(b) Accept rate under different B^{ex}

Figure 10: Accept rate under different B^{in} and B^{ex}

middleboxes exist, we then check the number of APPs in the two subtrees and
400 exchange the smaller one with the MBs to increase the degree of colocation for
APPs. We excute the exchanging until all exchangeable MBs or APPs have
been exchanged. If current exchangeable MBs have been exchanged completely,
we then select the next exchangeable type of MBs to exchange (line 9). After
each exchanging, we update the number of MBs and APPs in the two sides.




(a) VM slot utilization under different B^{in}

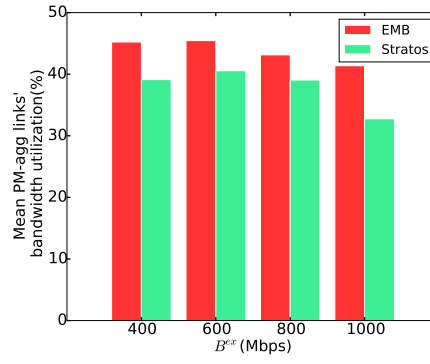


(b) VM slot utilization under different B^{ex}

Figure 11: VM slot utilization under different B^{in} and B^{ex}



(a) Core-aggregation links' bandwidth utilization



(b) ToR-PM links' bandwidth utilization

Figure 12: Bandwidth utilization under different B^{ex}

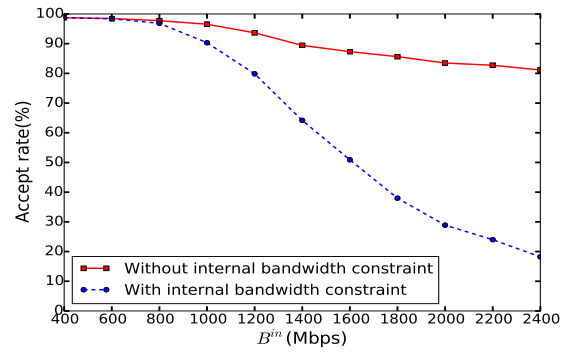


Figure 13: Accept rate with or without internal bandwidth constraint

405 5. Evaluation

We evaluate our placement algorithm via simulation. We build a tree-like three-layer network topology, which is inspired by a real cloud datacenter with 30000+ VMs. Each ToR switch is connected to 40 physical servers with 10Gbps NICs. And every PM is divided into 8 slots. The oversubscription ration are
 410 4:1 at the leaf level and 10:1 at the spine layer.

To simulate the tenant request, we model the intra-tenant bandwidth B^{in} and inter-tenant bandwidth B^{ex} to follow normal distribution. The maximal capacity b_k of each type of middleboxes also follows the normal distribution with average value as 1000. For simplicity, we assume that each tenant request only
 415 contains one middlebox chain. The length of the middlebox chain is randomly selected in [1,3]. And the number of APP VMs required corresponding to a middlebox chain is equally distributed between 3 to 10. The fraction of open tenants is set to be 10% while the fraction of client tenants is defined as 25%. Furthermore, there are also some tenants that have communication request with
 420 each other. We set the fraction of these tenants to be 15%.

We consider two placement algorithms: 1) a VM placement strategy (without scaling requirement) described in Stratos [2] which logically partitions the VMs into per-rack partitions and places them with minimal inter-partition communication; 2) our placement algorithm(EMB) that places the MB and APP
 425 VMs carefully given in Algorithm 1. To compare the two placement algorithms, we create tenant requests continuously until the datacenter could not accept tenant request any more.

We first evaluate the tenant accept rates under different B^{in} and B^{ex} in Fig. 10. In Fig. 10(a), we consider the accept rate of tenant requests with various
 430 intra-tenant communication demands (B^{in}). $B^{in}=1200$ means that the tenants B^{in} follows the normal distribution with average value as 1200 and variation equal to $B^{in}/3=400$. In Fig. 10(a), we find that both algorithms have similar accept rates with varying B^{in} .

The results of accept rate with varying external bandwidth (B^{ex}) are shown

435 in Fig. 10(b). Similarly, $B^{ex}=1200$ means that B^{ex} follows the normal distribu-
 tion with average value of 1200 and variation as $B^{ex}/3=400$. Obviously, EMB
 can always accept tenant requests reposefully. And even with large B^{ex} , EMB
 can keep a relatively high accept rate. While the accept rate of Stratos drops
 quickly when B^{ex} is greater than 800Mbps. This is because Stratos usually
 440 occupies more bandwidth resource in higher level links. More specifically, the
 number of total VMs required for tenant requests increases as B^{ex} increases
 since we need more middlebox instances to process a larger external traffic.
 This leads to an increased demand for VM slots resource in the data center,
 and the tenant requests are forced to spread their VMs across more racks. The
 445 placement strategy used in Stratos is more likely to select the racks that reside
 across several subtrees to accommodate the tenant request, which may cause a
 big waste of bandwidth in the high-level links. While our placement algorithm
 always tries to select the lowest-level subtree that has sufficient VM slots for
 the tenant, which can save the precious bandwidth resource of the high level
 450 links effectively. Thus, EMB is helpful to allow the provider to accommodate
 more tenant requests while reducing the bandwidth utilization of high-level links
 which allows more future tenant requests to be accepted.

Fig. 11 shows the VM slots utilization of the datacenter under different B^{in}
 and B^{ex} . We can see that the results of VM slots utilization have similar chang-
 455 ing trends compared to the accept rates shown in Fig. 10. This is reasonable
 since higher accept rates represent that there are more requests accepted and
 more VM slots resources allocated.

For a better understanding of above results, we evaluate the downlink band-
 width utilization of core-aggregation and ToR-PM links. Note that the servers
 460 are on the same layer of the three-level network topology, if there is one download
 bandwidth consumption, there must be another upload bandwidth consumption
 correspondingly. This leads to the result that the download and upload results
 are similar to each other. Hence, we just show the results of the download band-
 width here. The results are shown in Fig. 12, where the x-coordinate shows
 465 the intra-tenant bandwidth (B^{in}) consumption, and the y-coordinate shows the

average bandwidth utilization of the core-aggregation and ToR-PM download links in the two subfigures respectively. In Fig. 12(a), we can see that EMB has lower bandwidth utilization than Stratos in the core-aggregation links. On the other side, EMB has a higher utilization of the PM-ToR links shown in Fig. 12(b) which, in turn, proves the effect of lowest-level subtrees locality placement of our algorithm. To sum up, EMB can save the bandwidth efficiently of high-level links which are usually the bottlenecks for most oversubscription data center networks while making the best of low-level links bandwidth.

We then evaluate the impact of B^{in} and bandwidth constraint inside a PM. The bandwidth constraint within a physical machine limits the number of VMs allowed to be placed. The experimental results are shown in Fig. 13. It is observable that when internal bandwidth constraint is considered, the accept rate of tenant requests quickly drops as B^{in} increases. As the link capacity and internal constraint are the same in our simulation, a request is always limited by the internal PM constraint first, which reduces the number of available VM slots and leads to a lower accept rate.

6. Conclusion

Since middleboxes are widely used in current networks, it is essential to consider middleboxes when embedding predicable virtual networks. This paper tries to fill the gap. To achieve this goal, this paper first models the tenant's requirement of software middleboxes and applications. By augmenting software middleboxes as VMs in the Tenant Application Graph, an VMs placement algorithm, called EMB, is proposed to place these VMs (including both normal application VMs and virtual middleboxes) in the underlying infrastructure. The evaluation results have shown that the proposed virtual middleboxes placement algorithm can effectively use the resource of multi-tenant datacenters and provide performance guarantee for tenants.

References

- [1] A. Gember, R. Grandl, J. Khalid, A. Akella, Design and implementation
495 of a framework for software-defined middlebox networking, in: ACM SIG-
COMM 2013, 2013, pp. 467–468.
- [2] A. Gember, A. Akella, A. Anand, T. Benson, R. Grandl, Stratos: Vir-
tual middleboxes as first-class entities, Tech. Rep. TR1771, University of
Wisconsin-Madison.
- 500 [3] H. Ballani, K. Jang, T. Karagiannis, C. Kim, D. Gunawardena, G. O’Shea,
Chatty tenants and the cloud network sharing problem., in: NSDI, 2013,
pp. 171–184.
- [4] J. Lee, Y. Turner, M. Lee, L. Popa, S. Banerjee, J.-M. Kang, P. Sharma,
Application-driven bandwidth guarantees in datacenters, in: ACM SIG-
505 COMM, 2014, pp. 467–478.
- [5] H. Ballani, P. Costa, T. Karagiannis, A. Rowstron, Towards predictable
datacenter networks, in: ACM SIGCOMM Computer Communication Re-
view, Vol. 41, 2011, pp. 242–253.
- [6] A. Williamson, Has amazon ec2 become over subscribed, online]
510 [http://alan.blog-city.com/has amazon ec2 become over subscribed.htm](http://alan.blog-city.com/has_amazon_ec2_become_over_subscribed.htm).
- [7] P. Patel, D. Bansal, L. Yuan, A. Murthy, A. Greenberg, D. A. Maltz,
R. Kern, H. Kumar, M. Zikos, H. Wu, et al., Ananta: cloud scale load
balancing, in: ACM SIGCOMM 2013, 2013, pp. 207–218.
- [8] J. Martins, M. Ahmed, C. Raiciu, V. Olteanu, M. Honda, R. Bifulco,
515 F. Huici, Clickos and the art of network function virtualization, in: NSDI,
2014, pp. 459–473.
- [9] R. Yu, G. Xue, V. Kilari, X. Zhang, Network function virtualization in the
multi-tenant cloud, IEEE Network 29 (3) (2015) 42–47.

- [10] I. web portal, Network functions virtualisation., in:
520 <http://portal.etsi.org/portal/server.pt/community/NFV/367>, Retrieved
20 June 2013.
- [11] M.Dobrescu, N.Egi, S.Ratnasamy, Toward predictable performance in soft-
ware packet-processing platforms, in: NSDI, 2012.
- [12] V. Sekar, N. Egi, S. Ratnasamy, M. K. Reiter, G. Shi, Design and imple-
525 mentation of a consolidated middlebox architecture., in: NSDI, 2012, pp.
323–336.
- [13] L. Popa, G. Kumar, M. Chowdhury, A. Krishnamurthy, S. Ratnasamy,
I. Stoica, Faircloud: sharing the network in cloud computing, in: the ACM
SIGCOMM 2012, 2012, pp. 187–198.
- 530 [14] R. Mijumbi, J. Serrat, J. Gorricho, N. Bouten, F. De Turck, R. Boutaba,
Network function virtualization: State-of-the-art and research challenges,
IEEE Communications Surveys Tutorials PP (99) (2015) 1–1.
- [15] Huawei nfV open lab, in: <http://pr.huawei.com/en/news>, 2015.
- [16] Cloudnfv, in: <http://www.cloudnfv.com/>, 2015.
- 535 [17] S. Mehraghdam, M. Keller, H. Karl, Specifying and placing chains of vir-
tual network functions, in: IEEE 3rd International Conference on Cloud
Networking (CloudNet), 2014.
- [18] R. Cohen, L. Lewin-Eytan, J. S. Naor, D. Raz, Near optimal placement of
virtual network functions, in: IEEE INFOCOM, 2015, pp. 1346–1354.
- 540 [19] C. Chekuri, S. Khanna, On multi-dimensional packing problems, in:
SODA, 1999.
- [20] A. Fischer, J. F. Botero, M. Till Beck, H. De Meer, X. Hesselbach, Virtual
network embedding: A survey, IEEE Communications Surveys & Tutorials
15 (4) (2013) 1888–1906.

- 545 [21] PCI-SIG SR-IOV, http://www.pcisig.com/iov/single_root/.
- [22] J. Hwang, K. Ramakrishnan, T. Wood, Netvm: high performance and flexible networking using virtualization on commodity platforms, in: NSDI, 2014.

Appendix

550 To address the problem of unpredictable MB performance, instead of allocating new resources of VMs and network after MBs are placed in datacenters, we introduce a probability model to represent the variable resource requirement of individual tenant. Through statistically multiplexing on VMs, we are able to balance the physical resources consumption and scalability requirement of software MBs.

We re-examine the MBs request and take the time-varying resource requirements of MBs into consideration. Specifically, the MBs' requirement is composed of a basic requirement and a variable requirement with probability. Based on this model, we replace the $\langle T, M \rangle$ in Section 3.2 with tuples $\langle T, M, m, P \rangle$, 560 where T and M denote the type of MBs and the corresponding number of VMs in the basic requirement, and m denotes the number of VMs in the variable subrequirement with probability P . For example, given $T = \{\text{FW}, \text{IDS}, \text{LB}\}$, $M = \{2, 4, 2\}$, $m = \{1, 2, 1\}$, $P = \{0.3, 0.2, 0.3\}$, we then know that the tenant request for 2 firewall MBs and the request may increase to $2 + 1 = 3$ MBs with probability of 0.2. Similarly, 4 IDS MBs (resp. 2 LB MBs) are needed with probability 0.8 (resp. 0.7) and 6 IDS MBs (resp. 3 LB MBs) with a probability of 0.2 (resp. 0.3). In summary, a tenant's request can be specified by multiple virtual networks (VNs) and each VN can be specified by an eight-tuple $\langle N, B^{in}, B^{ex}, dependencies, T, M, m, P \rangle$. It allows elasticity by having 570 a probability model for the number of MBs needed in each VN.

If multiple units of subrequirement do not occur simultaneously, we can provide few redundancy to tradeoff between utilization and availability. By

sharing free VMs between variable subrequirements, a specific MB can be soon deployed on the free VM once the scale up/out requirement comes. Although the subrequirement occurs with a probability less than 1, the collision may happen if there is not enough VM provided. To guarantee the performance, we apply the probabilistic model to solve this problem. For simplicity, we reuse the notation M to indicate how many types of MBs are required by tenant. Let X_i indicate whether the i -th subrequirement occurs. Assuming the resource requirement of different middleboxes are independent, the probability of the collision happening on one supplement VM is

$$\begin{aligned} & Pr \left[\sum_{i=1}^M X_i > 1 \right] \\ &= 1 - \prod_{i=1}^M (1 - p_i) - \sum_{i=1}^M \left(p_i \prod_{k \in [1, M], k \neq i} (1 - p_k) \right) \end{aligned} \quad (.1)$$

By providing enough number of free VMs during virtual middleboxes network deployment, the probability of collision can be guaranteed less than a collision threshold ρ when the system is on board. The number of subrequirement v can be determined by solving the following function:

$$Pr \left[\sum_{i=1}^M X_i \leq v \right] \leq \rho \quad (.2)$$

The collision probability $Pr \left[\sum_{i=1}^M X_i \leq v \right]$ can be achieved based on the including excluding principle. Therefore, the total number of MB VMs to be allocated for one tenant is $|M| + v$, where $|M|$ is the number of basic MBs requirement and v is the additional free VMs to be allocated as backup resources. Therefore,

575 Fig. .14 (B) uses the cascade of TAG model to abstract the communication dependency of VNs given in Fig. .14 (A). In comparison with Fig. 4, the VMs inside the blue dotted line are pre-allocated VMs for middleboxes subrequirement, which are shared by both types of middleboxes requests. If there are scale-out or failover request, specified MBs will be placed on these VMs
580 to provide performance guarantee. If a tenant requires for multiple types of MBs, we can place free slots between cascade two types of MBs according to

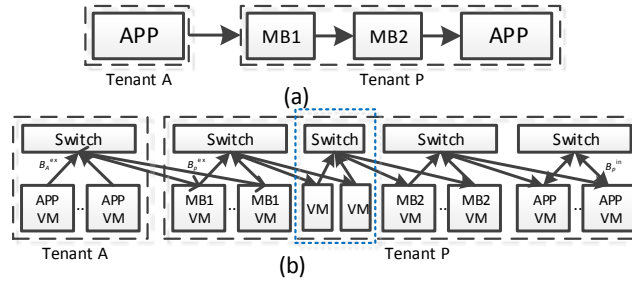


Figure .14: Virtual Network Model. (a) the communication between two tenants; (b) cascade of TAG

the expectation of their scalability.